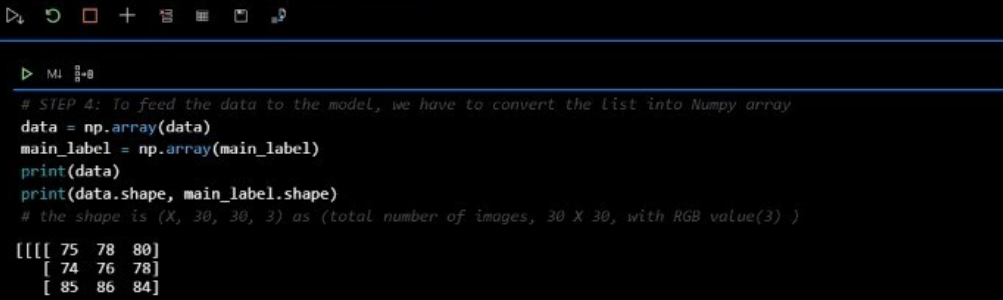


OUTPUTS

Output 1:

Data which is converted into a numpy array i.e. data and labels



The screenshot shows the Visual Studio Code interface with the 'main.ipynb' file open. The notebook cell [8] contains the following code:

```
[8]: > MI 0=0

# STEP 4: To feed the data to the model, we have to convert the list into Numpy array
data = np.array(data)
main_label = np.array(main_label)
print(data)
print(data.shape, main_label.shape)
# the shape is (X, 30, 30, 3) as (total number of images, 30 X 30, with RGB value(3) )

[[[ 75  78  80]
 [ 74  76  78]
 [ 85  86  84]
 ...
 [ 68  75  74]
 [ 65  69  68]
 [ 66  67  66]]

 [[ 83  84  86]
 [ 80  80  82]
 [ 88  88  83]
 ...
 [ 73  77  78]
 [ 76  78  75]
 [ 80  80  78]]

 [[ 78  78  80]
 [ 86  85  86]
 [ 90  89  90]
 ...
 [ 71  74  71]
 [ 73  74  69]
 [ 78  78  74]]

...
```

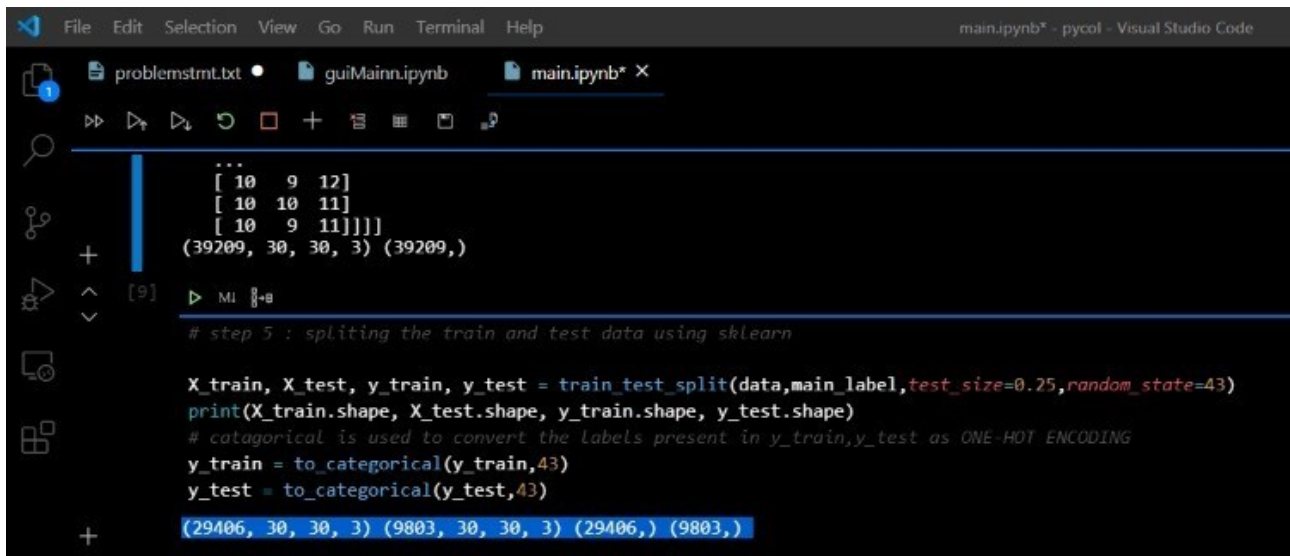
Output 2:

Shape of the dataset -> the list which contains the data and the labels that specifies the total number of images in the format (a,b,c,d) ->(shape of data,30x30 pixels, RGB colour)

[illegible]

Output 3:

Shape of training data and testing data, split in the ratio of 75:25 with random state of 43



```
File Edit Selection View Go Run Terminal Help main.ipynb* - pycol - Visual Studio Code

problemstmt.txt • guiMainn.ipynb main.ipynb x

[39209, 30, 30, 3] (39209,)

[9] In [ ]:

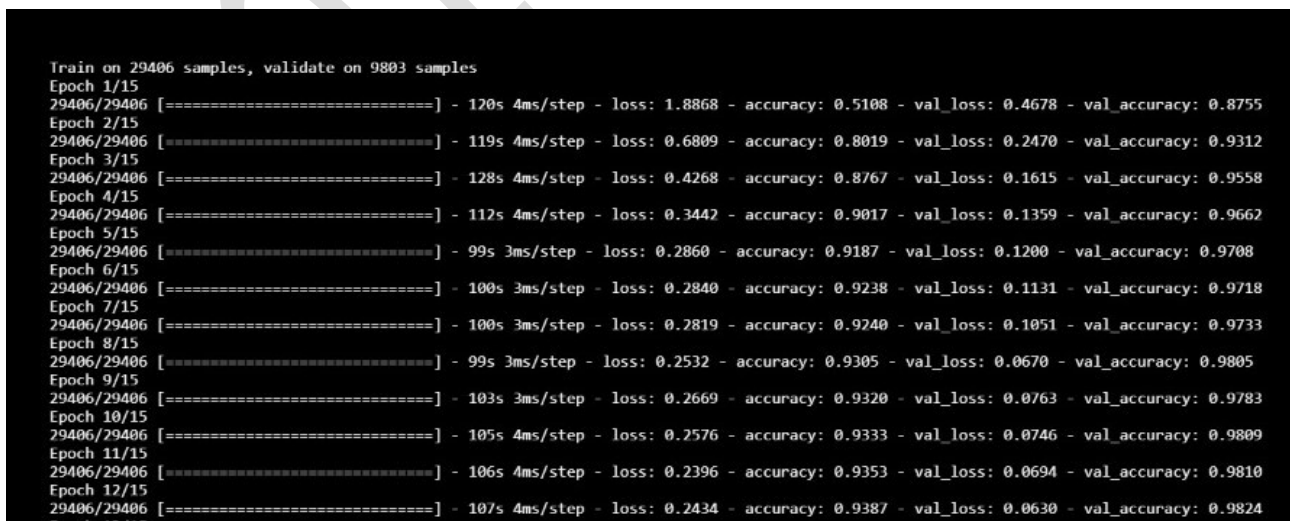
# step 5 : splitting the train and test data using sklearn

X_train, X_test, y_train, y_test = train_test_split(data,main_label,test_size=0.25,random_state=43)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
# catagorical is used to convert the labels present in y_train,y_test as ONE-HOT ENCODING
y_train = to_categorical(y_train,43)
y_test = to_categorical(y_test,43)

(29406, 30, 30, 3) (9803, 30, 30, 3) (29406,) (9803,)
```

Output 4:

fitting the training dataset using Keras



```
Train on 29406 samples, validate on 9803 samples
Epoch 1/15
29406/29406 [=====] - 120s 4ms/step - loss: 1.8868 - accuracy: 0.5108 - val_loss: 0.4678 - val_accuracy: 0.8755
Epoch 2/15
29406/29406 [=====] - 119s 4ms/step - loss: 0.6809 - accuracy: 0.8019 - val_loss: 0.2470 - val_accuracy: 0.9312
Epoch 3/15
29406/29406 [=====] - 128s 4ms/step - loss: 0.4268 - accuracy: 0.8767 - val_loss: 0.1615 - val_accuracy: 0.9558
Epoch 4/15
29406/29406 [=====] - 112s 4ms/step - loss: 0.3442 - accuracy: 0.9017 - val_loss: 0.1359 - val_accuracy: 0.9662
Epoch 5/15
29406/29406 [=====] - 99s 3ms/step - loss: 0.2860 - accuracy: 0.9187 - val_loss: 0.1200 - val_accuracy: 0.9708
Epoch 6/15
29406/29406 [=====] - 100s 3ms/step - loss: 0.2840 - accuracy: 0.9238 - val_loss: 0.1131 - val_accuracy: 0.9718
Epoch 7/15
29406/29406 [=====] - 100s 3ms/step - loss: 0.2819 - accuracy: 0.9240 - val_loss: 0.1051 - val_accuracy: 0.9733
Epoch 8/15
29406/29406 [=====] - 99s 3ms/step - loss: 0.2532 - accuracy: 0.9305 - val_loss: 0.0670 - val_accuracy: 0.9805
Epoch 9/15
29406/29406 [=====] - 103s 3ms/step - loss: 0.2669 - accuracy: 0.9320 - val_loss: 0.0763 - val_accuracy: 0.9783
Epoch 10/15
29406/29406 [=====] - 105s 4ms/step - loss: 0.2576 - accuracy: 0.9333 - val_loss: 0.0746 - val_accuracy: 0.9809
Epoch 11/15
29406/29406 [=====] - 106s 4ms/step - loss: 0.2396 - accuracy: 0.9353 - val_loss: 0.0694 - val_accuracy: 0.9810
Epoch 12/15
29406/29406 [=====] - 107s 4ms/step - loss: 0.2434 - accuracy: 0.9387 - val_loss: 0.0630 - val_accuracy: 0.9824
```

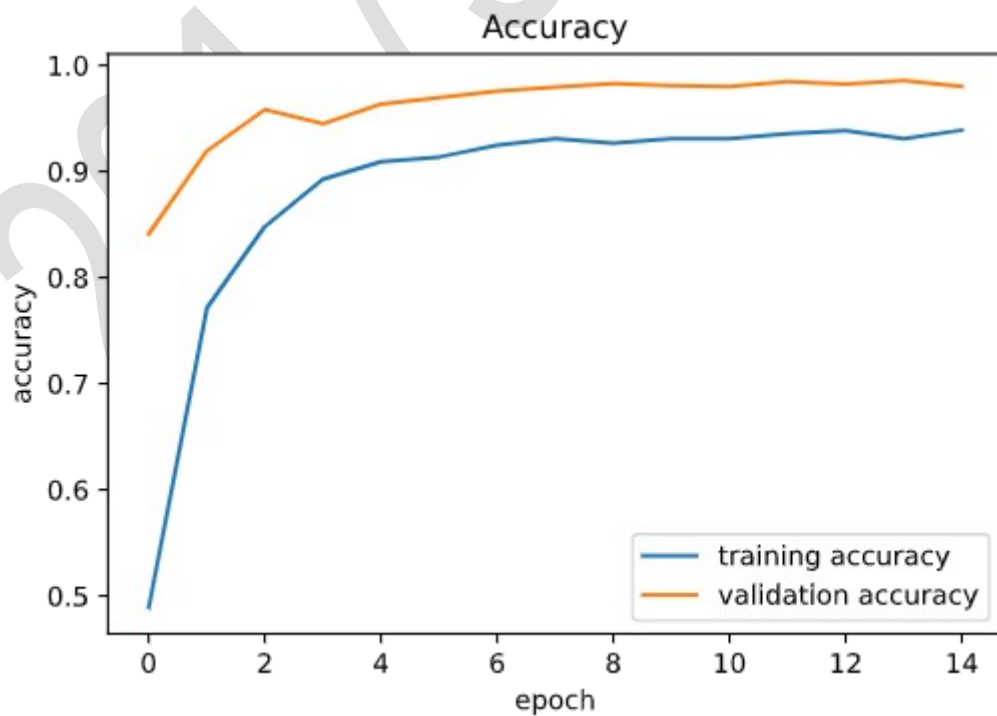
```

29406/29406 [=====] - 128s 4ms/step - loss: 0.4268 - accuracy: 0.8767 - val_loss: 0.1615 - val_accuracy: 0.9558
Epoch 4/15
29406/29406 [=====] - 112s 4ms/step - loss: 0.3442 - accuracy: 0.9017 - val_loss: 0.1359 - val_accuracy: 0.9662
Epoch 5/15
29406/29406 [=====] - 99s 3ms/step - loss: 0.2860 - accuracy: 0.9187 - val_loss: 0.1200 - val_accuracy: 0.9708
Epoch 6/15
29406/29406 [=====] - 100s 3ms/step - loss: 0.2840 - accuracy: 0.9238 - val_loss: 0.1131 - val_accuracy: 0.9718
Epoch 7/15
29406/29406 [=====] - 100s 3ms/step - loss: 0.2819 - accuracy: 0.9240 - val_loss: 0.1051 - val_accuracy: 0.9733
Epoch 8/15
29406/29406 [=====] - 99s 3ms/step - loss: 0.2532 - accuracy: 0.9305 - val_loss: 0.0670 - val_accuracy: 0.9805
Epoch 9/15
29406/29406 [=====] - 103s 3ms/step - loss: 0.2669 - accuracy: 0.9320 - val_loss: 0.0763 - val_accuracy: 0.9783
Epoch 10/15
29406/29406 [=====] - 105s 4ms/step - loss: 0.2576 - accuracy: 0.9333 - val_loss: 0.0746 - val_accuracy: 0.9809
Epoch 11/15
29406/29406 [=====] - 106s 4ms/step - loss: 0.2396 - accuracy: 0.9353 - val_loss: 0.0694 - val_accuracy: 0.9810
Epoch 12/15
29406/29406 [=====] - 107s 4ms/step - loss: 0.2434 - accuracy: 0.9387 - val_loss: 0.0630 - val_accuracy: 0.9824
Epoch 13/15
29406/29406 [=====] - 107s 4ms/step - loss: 0.2294 - accuracy: 0.9437 - val_loss: 0.0603 - val_accuracy: 0.9848
Epoch 14/15
29406/29406 [=====] - 108s 4ms/step - loss: 0.2151 - accuracy: 0.9459 - val_loss: 0.0512 - val_accuracy: 0.9867
Epoch 15/15
29406/29406 [=====] - 73s 2ms/step - loss: 0.2195 - accuracy: 0.9447 - val_loss: 0.0801 - val_accuracy: 0.9791

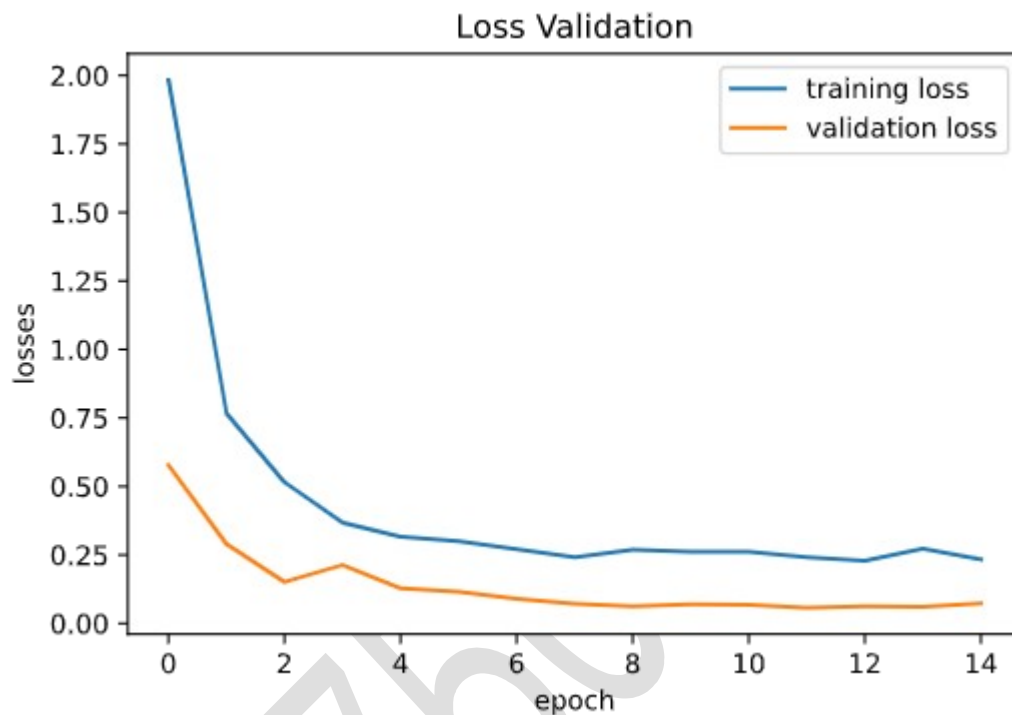
```

Output 5:

Accuracy Graph



Loss Graph



Output 6:

head of the test CSV dataset

```
File Edit Selection View Go Run Terminal Help
main.ipynb - pycol - Visual Studio Code

problemstmt.txt • guiMainn.ipynb main.ipynb ×

[14] In [14]: #STEP 9: The test.csv has the details on the image, path and their respective labels
y_test = pd.read_csv('Dataset\Test.csv')
#debugging
print(y_test.head())
```

	Width	Height	Roi.X1	Roi.Y1	Roi.X2	Roi.Y2	ClassId	Path
0	53	54	6	5	48	49	16	Test/00000.png
1	42	45	5	5	36	40	1	Test/00001.png
2	48	52	6	6	43	47	38	Test/00002.png
3	27	29	5	5	22	24	33	Test/00003.png
4	60	57	5	5	55	52	11	Test/00004.png

Output 7:

accuracy

```
#accuracy
print(accuracy_score(labels, pred))
+ 0.9468725257323832
```

Output 8:

classification report

```
print(classification_report(labels,pred))
precision    recall  f1-score   support

   0:  1.00   0.98   0.99     60
   1:  0.99   0.99   0.99    720
   2:  0.99   0.97   0.98    750
   3:  0.91   0.88   0.90    450
   4:  0.99   0.96   0.98    660
   5:  0.81   0.97   0.88    630
   6:  0.94   0.79   0.86    150
   7:  0.95   0.94   0.95    450
   8:  0.97   0.92   0.94    450
   9:  0.99   0.97   0.98    480
  10:  0.97   0.97   0.97    660
  11:  0.86   0.91   0.88    420
  12:  0.98   0.93   0.96    690
  13:  0.99   1.00   0.99    720
  14:  0.99   0.98   0.99    270
  15:  0.91   0.99   0.95    210
  16:  0.89   0.99   0.93    150
  17:  0.96   0.99   0.98    360
  18:  0.98   0.87   0.92    390
  19:  0.97   0.93   0.95     60
  20:  0.96   0.94   0.95     90
  21:  0.95   0.90   0.93     90
  22:  1.00   0.96   0.98    120
-- -- -- -- --
```

```
print(classification_report(labels,pred))
-- -- -- -- --
 23:  0.96   0.92   0.94    150
 24:  0.99   0.88   0.93     90
 25:  0.97   0.91   0.94    480
 26:  0.87   0.84   0.86    180
 27:  0.97   0.48   0.64     60
 28:  0.98   0.93   0.96    150
 29:  0.92   0.91   0.92     90
 30:  0.71   0.80   0.75    150
 31:  0.86   0.99   0.92    270
 32:  0.88   0.97   0.92     60
 33:  0.97   0.90   0.98    210
 34:  0.98   0.99   0.99    120
 35:  0.99   0.97   0.98    390
 36:  0.98   0.97   0.98    120
 37:  0.97   0.98   0.98     60
 38:  0.95   0.99   0.97    690
 39:  0.94   0.83   0.88     90
 40:  0.98   0.98   0.98     90
 41:  0.91   0.72   0.80     60
 42:  0.80   0.96   0.87     90

 accuracy          0.95  12630
 macro avg         0.94   0.92   0.93  12630
 weighted avg      0.95   0.95   0.95  12630
```

Final accuracy: 94%

Final Output:

